


Sorting Exercise

Jonas Schenke
28.01.2019

Concept

1.Stream data in:



2.Synchronize threads

3.Stream data out:



4.Flush file

Current Implementation

- Implementation in C++
- Everything loaded to RAM (has to fit!)
- Each Quicksort executed in separate thread
- Select smallest element from lists
 - create output blocks for streaming
- Single threaded merging

Configuration

- Block size ~ 610MiBi
 - Input file size ~1.9GiBi (20 mio. Entries)
 - Test System:
 - Ryzen 1700 (16*3.2GHz)
 - 16 GiBi RAM
 - 240GB Corsair MP500 NVMe SSD (read:~3GB/s)
- => ~10s runtime / 2 mio. entries/s**

Further Notes

- Time for input * 2 = (ca.) time for output
- ~140% CPU usage
 - Input multithreaded
 - Output single threaded → improvable
- Potential speedup with parallel input sorting algorithm
- Merging algorithm improvable
- Code adaptable using template metaprogramming